

Introduction to Information Retrieval

<http://informationretrieval.org>

IIR 7: Scores in a Complete Search System

Mihai Surdeanu

(Based on slides by Hinrich Schütze at informationretrieval.org)

Spring 2017

Overview

- 1 Why rank?
- 2 The complete search system
- 3 Implementation of ranking

Take-away today

- The importance of ranking: User studies at Google
- The complete search system
- Implementation of ranking

Outline


- 1 Why rank?
- 2 The complete search system
- 3 Implementation of ranking

Why is ranking so important?

- Last lecture: Problems with unranked retrieval
 - Users want to look at a few results – not thousands.
 - It's very hard to write queries that produce a few results.
 - Even for expert searchers
 - → Ranking is important because it effectively **reduces a large set of results to a very small one.**
- Next: More data on “users only look at a few results”
- Actually, in the vast majority of cases they only examine 1, 2, or 3 results.

Empirical investigation of the effect of ranking

- The following slides are from Dan Russell's JCDL talk
- Dan Russell was the “Über Tech Lead for Search Quality & User Happiness” at Google.
- How can we measure how important ranking is?
- Observe what searchers do when they are searching in a controlled setting
 - Videotape them
 - Ask them to “think aloud”
 - Interview them
 - Eye-track them
 - Time them
 - Record and count their clicks



So.. Did you notice the FTD official site?

To be honest, I didn't even look at that.

At first I saw "from \$20" and \$20 is what I was looking for.

To be honest, 1800-flowers is what I'm familiar with and why I went there next even though I kind of assumed they wouldn't have \$20 flowers

And you knew they were expensive?

I knew they were expensive but I thought "hey, maybe they've got some flowers for under \$20 here..."

But you didn't notice the FTD?

No I didn't, actually... that's really funny.

Interview video

Rapidly scanning the results

Note scan pattern:

Page 3:

Result 1
Result 2
Result 3
Result 4
Result 3
Result 2
Result 4
Result 5
Result 6 <click>

Q: Why do this?

A: What's learned later influences judgment of earlier content.

The screenshot shows a Google search for "children's unicycle". The results are numbered 1 through 6, with red arrows indicating a scanning pattern that starts at the top, moves down to result 1, then up to result 2, then down to result 3, then up to result 4, then down to result 5, and finally up to result 6. The search results are as follows:

- 1 **Unicycle UK.com - F.A.Q. - What size?**
12" wheel **unicycle**: this is a small **children's unicycle** size. It's good for **children** who are too small to ride a 16" **unicycle**, but it needs smooth ground ...
www.unicycle.uk.com/FAQ.asp?Category=53 - 23k - Cached - Similar pages
- 2 **Selecting a unicycle: Unicycle.com NZ : buy a unicycle or learn ...**
16" wheel **unicycle**: this is a **children's unicycle**, the small wheel makes it only suitable for smooth areas. Best used indoors or on smooth ground; ...
www.unicycle.co.nz/View.php?action=Page&Name=Selectingaunicycle - 22k - Cached - Similar pages
- 3 **100 Miles for Kids - The Goal**
The Afghan Mobile Mini Circus for **Children** is an established ... attempt to break the GUINNESS WORLD RECORD for the ONE HOUR **UNICYCLE** DISTANCE RECORD. ...
www.unicycle4kids.org/ - 9k - Cached - Similar pages
- 4 **Unicycles page at Juggling World**
This is a **children's unicycle**, the small wheel makes it only suitable for very smooth areas. Best used indoors or on smooth ground, not so good outdoors ...
www.jugglingworld.biz/shop/products_unicycles.html - 100k - Cached - Similar pages
- 5 **Buy a Unicycle: Unicycle.com AU : buy a unicycle or learn unicycling**
Check out a Unicycle Learners Pack for an easy and economical way to take your first steps into the One Wheeled World ... Suitable as a **Children's Unicycle**. ...
www.unicycle.au.com/View.php?action=Page&Name=Unicycles - 10k - Cached - Similar pages
- 6 **Article - News - A unicycle ride for children**
Adam Brody, 21, of San Juan Capistrano, led a charity event Saturday that benefits the Orangewood **Children's** Foundation. The **Unicycle** Club of Southern ...
www.ocregister.com/ocregister/news/homepage/article_1293785.php - 31k - Cached - Similar pages

Kinds of behaviors we see in the data

Short / Nav




Topic exploration



Topic switch




Methodical results exploration




Query reform




Multitasking



Stacking behavior




Looking vs. Clicking



- Users view results one and two more often / thoroughly
- Users click most frequently on result one

Presentation bias – reversed results

- Order of presentation influences where users look **AND** where they click



Importance of ranking: Summary

- **Viewing abstracts:** Users are a lot more likely to read the abstracts of the top-ranked pages (1, 2, 3, 4) than the abstracts of the lower ranked pages (7, 8, 9, 10).
- **Clicking:** Distribution is even more skewed for clicking
- In 1 out of 2 cases, users click on the top-ranked page.
- Even if the top-ranked page is not relevant, 30% of users will click on it.
- → Getting the ranking right is very important.
- → Getting the top-ranked page right is most important.


Exercise

- Ranking is also one of the high barriers to entry for competitors to established players in the search engine market.
- Why?

Outline

- 1 Why rank?
- 2 The complete search system
- 3 Implementation of ranking

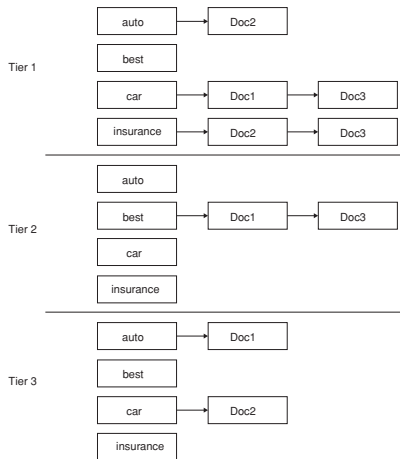
Complete search system



Tiered indexes

- Basic idea:
 - Create several tiers of indexes, corresponding to importance of indexing terms
 - During query processing, start with highest-tier index
 - If highest-tier index returns at least k (e.g., $k = 100$) results: stop and return results to user
 - If we've only found $< k$ hits: repeat for next index in tier cascade
- Example: two-tier system
 - Tier 1: Index of all titles
 - Tier 2: Index of the rest of documents
 - Pages containing the search words in the title are better hits than pages containing the search words in the body of the text.


Tiered index: by term frequency



Tiered indexes

- The use of tiered indexes is believed to be one of the reasons that Google search quality was significantly higher initially (2000/01) than that of competitors.
- (along with PageRank, use of anchor text and proximity constraints)
 - What are these?

Complete search system



Components we have introduced thus far

- Document preprocessing (linguistic and otherwise)
- Positional indexes
- Tiered indexes
- Spelling correction
- k-gram indexes for wildcard queries and spelling correction
- Query processing
- Document scoring

Components we haven't covered yet

- Document cache: we need this for generating snippets (= dynamic summaries; see Lecture 8)
- Zone indexes: They separate the indexes for different zones: the body of the document, all highlighted text in the document, anchor text, text in metadata fields etc. Why?
- Proximity ranking (e.g., rank documents in which the query terms occur in the same local window higher than documents in which the query terms occur far from each other)
- The two issues above best covered using machine-learned (ML) ranking functions!
 - Google didn't use ML for ranking until 2016. Why?
- Query parser

Tiered indexes vs. zone indexes

- Tiered index: partitions the *collection of documents*
- Zone index: partitions *individual documents*
- What does Lucene support?

Components we haven't covered yet: Query parser

- IR systems often guess what the user intended.
- The two-term query *London tower* (without quotes) may be interpreted as the phrase query *"London tower"*. **How?**
- The query *100 Madison Avenue, New York* may be interpreted as a request for a map.
- How do we “parse” the query and translate it into a formal specification containing phrase operators, proximity operators, indexes to search etc.? Need query syntax language!
 - We have seen Lucene's, but actual search engines do a lot of natural language processing as well, e.g., to recognize addresses, definitional questions, etc.

Exercise: Interactions with vector space retrieval

- How do we combine phrase retrieval with vector space retrieval?
- How do we combine Boolean retrieval with vector space retrieval?
- How do we combine wild cards with vector space retrieval?
- What does Lucene implement?

Exercise: Better tiered system

- Design criteria for tiered system
 - Each tier should be an order of magnitude smaller than the next tier.
 - Roughly: the top 100 hits for most queries should be in tier 1, the top 100 hits for most of the remaining queries in tier 2 etc.
 - We need a simple test for “can I stop at this tier or do I have to go to the next one?”
 - There is no advantage to tiering if we have to hit most tiers for most queries anyway.
- Question: Can you think of a better way of setting up a multitier system? Which “zones” of a document should be indexed in the different tiers (title, body of document, others?)? What criterion do you want to use for including a document in tier 1?

Outline

- 1 Why rank?
- 2 The complete search system
- 3 Implementation of ranking

Now we also need term frequencies in the index

Now we also need term frequencies in the index

BRUTUS	→	1,2	7,3	83,1	87,2	...
--------	---	-----	-----	------	------	-----

CAESAR	→	1,1	5,1	13,1	17,1	...
--------	---	-----	-----	------	------	-----

CALPURNIA	→	7,1	8,2	40,1	97,3
-----------	---	-----	-----	------	------

Now we also need term frequencies in the index

BRUTUS	→	1,2	7,3	83,1	87,2	...
--------	---	-----	-----	------	------	-----

CAESAR	→	1,1	5,1	13,1	17,1	...
--------	---	-----	-----	------	------	-----

CALPURNIA	→	7,1	8,2	40,1	97,3
-----------	---	-----	-----	------	------

term frequencies

Now we also need term frequencies in the index

BRUTUS	→	1,2	7,3	83,1	87,2	...
--------	---	-----	-----	------	------	-----

CAESAR	→	1,1	5,1	13,1	17,1	...
--------	---	-----	-----	------	------	-----

CALPURNIA	→	7,1	8,2	40,1	97,3
-----------	---	-----	-----	------	------

term frequencies

We also need positions. Not shown here.

Term frequencies in the inverted index

- Thus: In each posting, store $tf_{t,d}$ in addition to docID d .
- As an integer frequency, not as a (log-)weighted real number
- ...

Term frequencies in the inverted index

- Thus: In each posting, store $tf_{t,d}$ in addition to docID d .
- As an integer frequency, not as a (log-)weighted real number
...
- ... because real numbers are difficult to compress.
- Overall, additional space requirements are small: a byte per posting or less

Computing cosine: compare against the Boolean intersection!

COSINESCORE(q)

```
1  float Scores[ $N$ ] = 0
2  float Length[ $N$ ]
3  for each query term  $t$ 
4  do calculate  $w_{t,q}$  and fetch postings list for  $t$ 
5      for each pair( $d, tf_{t,d}$ ) in postings list
6      do Scores[ $d$ ] +  $= w_{t,d} \times w_{t,q}$ 
7  Read the array Length
8  for each  $d$ 
9  do Scores[ $d$ ] = Scores[ $d$ ] / Length[ $d$ ]
10 return Top  $k$  components of Scores[]
```

How do we compute the top k in ranking?

- We usually don't need a complete ranking.
- We just need the top k for a small k (e.g., $k = 100$).
- If we don't need a complete ranking, is there an efficient way of computing just the top k ?
- Naive:
 - Compute scores for all N documents
 - Sort
 - Return the top k
- Not very efficient
- Alternative: min heap

Use min heap for selecting top k out of N


- A binary min heap is a binary tree in which each node's value is less than the values of its children.
- It is a complete tree: all levels are completely filled, except possibly the last one. If the last level is not complete, the leaves are filled from left to right.
- Takes $O(N \log k)$ operations to construct (where N is the number of documents) ...
- ... then read off k winners in $O(k \log k)$ steps

Inserting into a min heap

- Place the new element in the next available position in the leaves.
- Compare the new element with its parent. If the new element is smaller, than swap it with its parent.
- Continue this process until either
 - the new elements parent is smaller than or equal to the new element, or
 - the new element reaches the root.

(Text adapted from CMU's Introduction to Data Structures)

Binary min heap: insert 0.75




Removing the smallest element from the min heap

- Place the root element in a variable to return later.
- Remove the last element in the deepest level and move it to the root.
- While the moved element has a value greater than at least one of its children, swap this value with the smaller-valued child.
- Return the original root that was saved.

(Text adapted from CMU's Introduction to Data Structures)

Binary min heap: remove 0.6



Selecting top k scoring documents in $O(N \log k)$

- Goal: Keep the top k documents seen so far
- Use a binary min heap
- To process a new document d' with score s' :
 - Get current minimum h_m of heap ($O(1)$)
 - If $s' \leq h_m$ skip to next document
 - If $s' > h_m$ heap-delete-root ($O(\log k)$)
 - Heap-add d'/s' ($O(\log k)$)

Even more efficient computation of top k ?

- Ranking has time complexity $O(N)$ where N is the number of documents.
- Optimizations reduce the constant factor, but they are still $O(N)$, $N > 10^{10}$
- Are there sublinear algorithms?

Even more efficient computation of top k ?

- Ranking has time complexity $O(N)$ where N is the number of documents.
- Optimizations reduce the constant factor, but they are still $O(N)$, $N > 10^{10}$
- Are there sublinear algorithms?
- What we're doing in effect: solving the k -nearest neighbor (kNN) problem for the query vector (= query point).
- There are no general solutions to this problem that are sublinear.

More efficient computation of top k : Three ideas (heuristics)

- Document-at-a-time processing
- Term-at-a-time processing
- Cluster pruning


Non-docID ordering of postings lists

- So far: postings lists have been ordered according to docID.
- Alternative: a query-independent measure of “goodness” of a page
- Example: **PageRank** $g(d)$ of page d , a measure of how many “good” pages hyperlink to d (chapter 21)
- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \text{cos}(q, d)$$

- This scheme supports early termination: We do not have to process postings lists in their entirety to find top k .

Non-docID ordering of postings lists (2)



► **Figure 7.2** A static quality-ordered index. In this example we assume that Doc1, Doc2 and Doc3 respectively have static quality scores $g(1) = 0.25$, $g(2) = 0.5$, $g(3) = 1$.

Postings no longer sorted by docID, but by g ! Does it matter?

Non-docID ordering of postings lists (3)

- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \cos(q, d)$$

Non-docID ordering of postings lists (3)

- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \cos(q, d)$$

- Suppose: (i) $g \rightarrow [0, 1]$; (ii) $g(d) < 0.1$ for the document d we're currently processing; (iii) smallest top k score we've found so far is 1.2

Non-docID ordering of postings lists (3)

- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \cos(q, d)$$

- Suppose: (i) $g \rightarrow [0, 1]$; (ii) $g(d) < 0.1$ for the document d we're currently processing; (iii) smallest top k score we've found so far is 1.2
- Then all subsequent scores will be < 1.1 .

Non-docID ordering of postings lists (3)

- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \cos(q, d)$$

- Suppose: (i) $g \rightarrow [0, 1]$; (ii) $g(d) < 0.1$ for the document d we're currently processing; (iii) smallest top k score we've found so far is 1.2
- Then all subsequent scores will be < 1.1 .
- So we've already found the top k and can stop processing the remainder of postings lists.

Non-docID ordering of postings lists (3)

- Order documents in postings lists according to PageRank:
 $g(d_1) > g(d_2) > g(d_3) > \dots$
- Define composite score of a document:

$$\text{net-score}(q, d) = g(d) + \cos(q, d)$$

- Suppose: (i) $g \rightarrow [0, 1]$; (ii) $g(d) < 0.1$ for the document d we're currently processing; (iii) smallest top k score we've found so far is 1.2
- Then all subsequent scores will be < 1.1 .
- So we've already found the top k and can stop processing the remainder of postings lists.
- Questions?


This was document-at-a-time processing

- Both docID-ordering and PageRank-ordering impose a consistent ordering on documents in postings lists.
- Computing cosines in this scheme is **document-at-a-time**
- This means we need to be careful to store intermediate results for all documents seen! (Remember the Scores[] variable in the cosine similarity algorithm?)
- Alternative: term-at-a-time processing

Term-at-a-time processing

- Idea 1 (sort postings):
 - Order the documents d in the postings list of term t by decreasing order of $tf_{t,d}$.
 - When traversing the postings list for a query term t , we stop either after a fixed number of documents r have been seen, or after the value of $tf_{t,d}$ has dropped below a threshold.
- Idea 2 (sort query):
 - Sort terms in query q in descending order of idf .
 - Stop after getting to terms with low idf values.

Cluster pruning



Cluster pruning: Algorithm

- At indexing time:
 - Pick \sqrt{N} documents at random from the collection as *leaders*.
 - For each document that is not a leader, compute its closest leader. These are *followers*.
- At search time:
 - Given a query q , find the leader L that is closest to q .
 - The candidate set consists of L together with all its followers. We compute the cosine scores *only* for the documents in this set.

Implementation of ranking: Summary

- Ranking is **very expensive** in applications where we have to compute similarity scores for all documents in the collection.
- In most applications, the vast majority of documents have **similarity score 0** for a given query → lots of potential for speeding things up.
- However, there is **no fast nearest neighbor algorithm** that is guaranteed to be correct even in this scenario.
- In practice: **use heuristics** to prune search space – usually works very well.

Take-away today

- The importance of ranking: User studies at Google
- The complete search system
- Implementation of ranking