

CSC 483/583: MIDTERM REVIEW OUTLINE

MIHAI SURDEANU

Before we begin:

- The exam will be open book. You are welcome to bring any text-books, notes, etc.
- You are allowed a simple, self-contained hand-calculator. Internet-connected devices are **not allowed** under any circumstances.

Topics to know for the midterm:

1. Lecture 1: Introduction and Boolean retrieval
 - A. Definition of information retrieval
 - B. Term-document incidence matrix: definition, how to build it
 - C. Inverted index: definition, how to build it, cost (runtime) of building it, why is it better than the incidence matrix?
 - D. Algorithm for intersection
 - E. Algorithm for other Boolean operators (see Homework #1)
 - F. Query optimization
2. Lecture 2: Term vocabulary and postings list
 - A. What is a document?
 - B. Token vs. term
 - C. Tokenization issues
 - D. Stop words, stemming, lemmatization
 - E. Skip pointers
 - F. Phrase queries, biword indexes
 - G. Positional indexes
 - H. Algorithm for proximity intersection
3. Lecture 3: Dictionaries and tolerant retrieval

Date: Fall 2015.

- A. Hashes vs. binary trees vs. B-trees
 - B. Permuterm trees
 - C. k-gram index
 - D. Edit distance, including reading out operations
 - E. Spelling correction using k-gram indexes
 - F. Context sensitive spelling correction
 - G. Soundex algorithm
- 4. Lecture 4: Index construction
 - A. Block sort-based indexing (BSBI)
 - B. Single-pass in-memory indexing (SPIMI)
 - C. Note: The remaining topics after SPIMI in this chapter are not required for the midterm
 - 5. Lecture 5: Index compression
 - A. Why compression?
 - B. Lossy vs. lossless compression
 - C. Heap's law
 - D. Zipf's law
 - E. Dictionary compression: as single string, string with blocking, front coding
 - F. Postings compression: gap encoding with variable-length encoding, gamma codes
 - 6. Lecture 6: vector space model
 - A. Feast of famine for Boolean queries
 - B. Jaccard coefficient: where else is this useful? Limitations
 - C. tf-idf
 - D. Vector space model
 - E. Cosine similarity
 - F. Different ways of encoding: term frequency, document frequency, normalization
 - 7. Lecture 7: Complete search system
 - A. User studies for ranking

- B. Tiered indexes
- C. Zone indexes, proximity ranking, scoring functions with multiple components
- D. Combinations of multiple scoring models, e.g., boolean and vector-space models, phrase-based and vector-space models
- E. Query parser
- F. Note: The remaining topics in this chapter (i.e., exact and inexact top K retrieval) are not required for the midterm