CSC 483/583: MIDTERM REVIEW OUTLINE

MIHAI SURDEANU

At the exam:

- You can bring a 2-page sheet with your notes.
- You are allowed a simple, self-contained hand-calculator. Internet-connected devices are **not allowed** under any circumstances.

Topics to know for the midterm:

1. Lecture 1: Introduction and Boolean retrieval
   A. Definition of information retrieval
   B. Term-document incidence matrix: definition, how to build it
   C. Inverted index: definition, how to build it, cost (runtime) of building it, why is it better than the incidence matrix?
   D. Algorithm for intersection
   E. Algorithms for other Boolean operators (see Homework #1)
   F. Query optimization
2. Lecture 2: Term vocabulary and postings list
   A. What is a document?
   B. Token vs. term
   C. Tokenization issues
   D. Stop words, stemming, lemmatization
   E. Skip pointers
   F. Phrase queries, biword indexes
   G. Positional indexes
   H. Algorithm for proximity intersection
3. Lecture 3: Dictionaries and tolerant retrieval
   A. Hashes vs. binary trees vs. B-trees

Date: Spring 2017.

   B. Permuterm trees
   C. k-gram index
   D. Edit distance, including reading out operations
   E. Spelling correction using k-gram indexes
   F. Context sensitive spelling correction
   G. Soundex algorithm
4. Lecture 4: Index construction
   A. Single-pass in-memory indexing (SPIMI)
   B. Note: Block sort-based indexing (BSBI), and the remaining topics after SPIMI in this chapter are not required for the midterm
5. Lecture 5: Index compression
   A. Why compression?
   B. Lossy vs. lossless compression
   C. Heap's law
   D. Zipf's law
   E. Dictionary compression not required for the midterm
   F. Postings compression: gap encoding with variable-length encoding, gamma codes
6. Lecture 6: vector space model
   A. Feast of famine for Boolean queries
   B. Jaccard coefficient: where else is this useful? Limitations
   C. tf-idf
   D. Vector space model
   E. Cosine similarity
   F. Different ways of encoding: term frequency, document frequency, normalization
7. Lecture 7: Complete search system
   A. User studies for ranking
   B. Tiered indexes
   C. Zone indexes, proximity ranking, scoring functions with multiple components

D. Combinations of multiple scoring models, e.g., boolean and vector-space models, phrase-based and vector-space models

E. Query parser

F. Exact top K retrieval using min heap

G. Inexact top K retrieval: document at a time, term at a time, cluster pruning

8. Lecture 8: Evaluation

A. Unranked evaluation: Precision, Recall, F score

B. Accuracy. Why is Accuracy not a good measure?

C. Ranked evaluation: P@1, precision-recall curve, mean average precision (MAP), mean reciprocal rank (MRR)

D. Inter-annotator agreement: Kappa measure

E. Real-world evaluations: A/B testing